

# Cloud Computing, Part 2

## Distributed and Pervasive Systems, MSc

Christian Marius Lillelund  
cl@ece.au.dk

Department of Electrical and Computer Engineering  
Aarhus University

Revised on February 5, 2021

# Outline

**Recap**

**Pods**

**Services**

**Deployments**

# Outline

## Recap

## Pods

## Services

## Deployments

# From last time

What did we learn?

- ▶ Cloud computing is delivery of online computing services
- ▶ Microservices are SoA-styled small, independent services that do one thing
- ▶ Containers are lightweight isolated software environments
- ▶ Docker is a container tool that can run and create containers
- ▶ Kubernetes is an open-source container orchestration engine

# Outline

Recap

**Pods**


Services

Deployments

# Pods in Kubernetes

What is a pod?

- ▶ A pod is a group of one or more tightly related containers that run together and share namespace
- ▶ Each pod is like a separate logical machine.
- ▶ All containers in a pod will appear to be running on the same logical machine.
- ▶ Can only run on one node



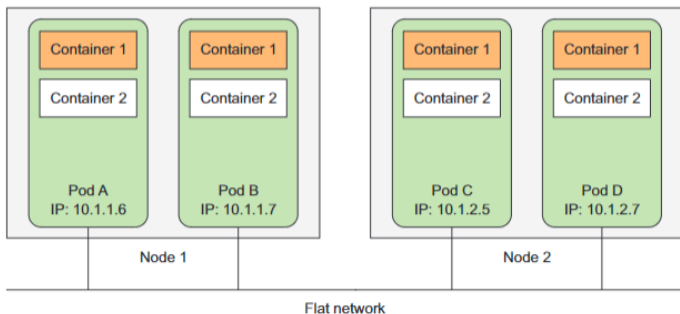
**Figure:** Fig. by courtesy of Marko Luksa[1]

# Why pods?

- ▶ Containers run only a single process.
- ▶ Pods allow us to bind containers together as a single unit.
- ▶ Pods run closely related processes together in the same environment.
- ▶ Processes think they are running together. Closed world.

## Network with pods

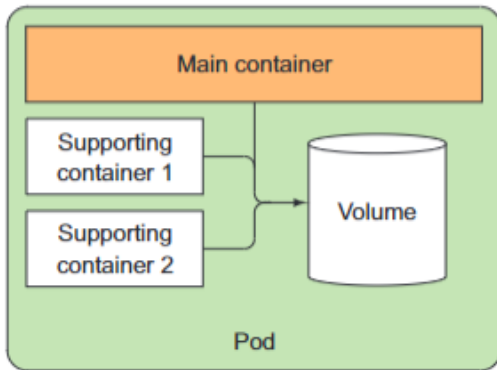
- ▶ All pods reside in a single flat, shared, network address space.
- ▶ Containers share the same IP
- ▶ Avoid port conflicts



**Figure:** Fig. by courtesy of Marko Luksa[1]



# The inside of a pod

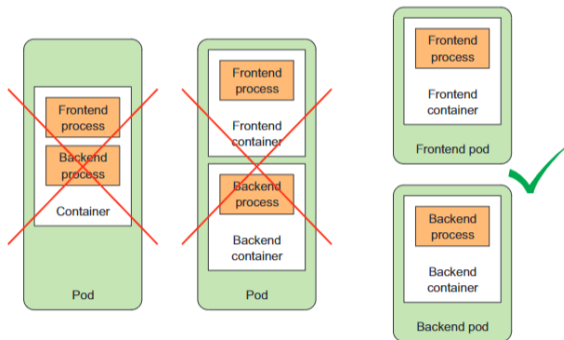


**Figure:** Fig. by courtesy of Marko Luksa[1]

## Using multiple containers

When to use multiple containers?

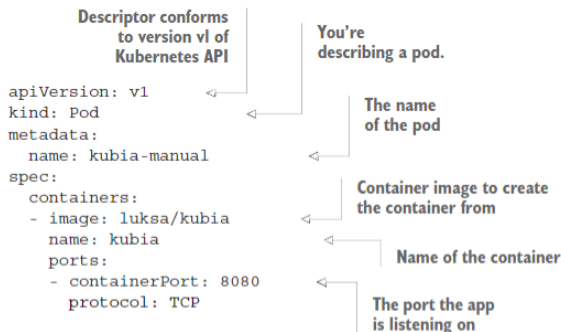
- ▶ Do they need to be run together?
- ▶ Do they scale together?
- ▶ Are they single components or one whole?



**Figure:** Fig. by courtesy of Marko Luksa[1]

# Creating pods

- ▶ Created by posting a YAML or JSON to the Kubernetes API
- ▶ Instead of "kubectl run", you post a YAML file
- ▶ Enables more options and version control



**Figure:** Listing by courtesy of Marko Luksa[1]

# Creating pods commands

Useful commands for creating pods and getting the manifest


```
$ kubectl create -f kubia-manual.yamlpod  
$ kubectl get po kubia-manual -o yaml  
$ kubectl get po kubia-manual -o json  
$ kubectl get pods  
$ kubectl logs kubia-manual
```

## Connecting to pods

Connect without a service

```
$ kubectl port-forward kubia-manual 8888:8080  
... Forwarding from 127.0.0.1:8888 -> 8080  
... Forwarding from [::1]:8888 -> 8080  
$ curl localhost:8888
```

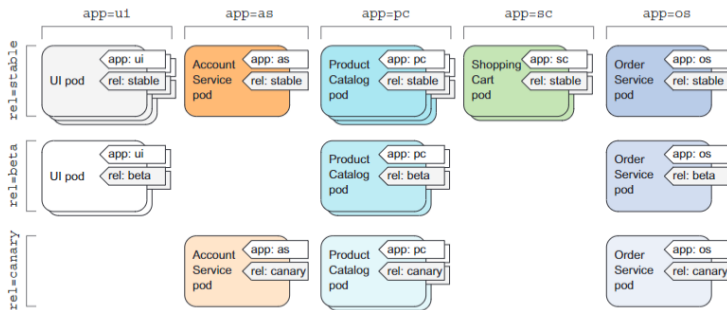
Note that minikube uses nodeIP:nodePort and not localhost:nodePort. Check with "minikube ip".



**Figure:** Fig. by courtesy of Marko Luksa[1]

# Organizing pods with labels

- ▶ Use labels to organize all Kubernetes resources.
- ▶ One or more labels
- ▶ Vertical and horizontal.



**Figure:** Fig. by courtesy of Marko Luksa[1]

# Organizing pods with labels cont.

```
apiVersion: v1
kind: Pod
metadata:
  name: kubia-manual-v2
  labels:
    creation_method: manual
    env: prod
spec:
  containers:
  - image: luksa/kubia
    name: kubia
    ports:
    - containerPort: 8080
      protocol: TCP
```

**Two labels are attached to the pod.**

**Figure:** Listing by courtesy of Marko Luksa[1]

## Organizing pods with labels cont.

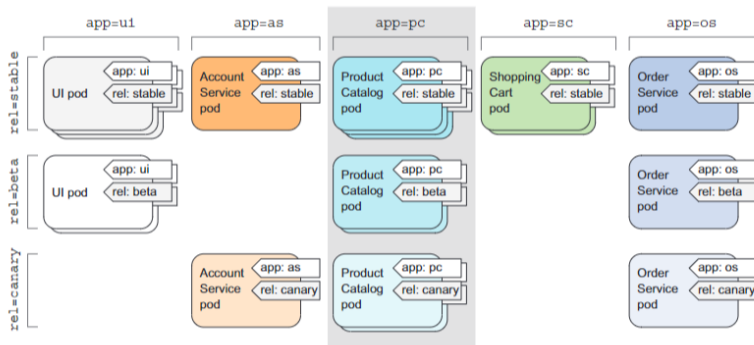
Create and show pods with labels

```
$ kubectl create -f kubia-manual-with-labels.yaml
$ kubectl get po --show-labels
$ kubectl get po -L creation_method,env
$ kubectl get po -l creation_method=manual
```

- ▶ Don't worry about scheduling. Kubernetes handles that.
- ▶ Never say specifically what node a pod should run on.



# Organizing pods with labels




**Figure:** Fig. by courtesy of Marko Luksa[1]

## Scheduling pods to specific nodes

Not best practice, but it is possible

```
$ kubectl label node gke-kubia-85f6-node-0rrx  
gpu=true  
$ kubectl get nodes -l gpu=true
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: kubia-gpu  
spec:  
  nodeSelector:  
    gpu: "true"  
  containers:  
  - image: luksa/kubia  
    name: kubia
```



nodeSelector tells Kubernetes to deploy this pod only to nodes containing the gpu=true label.

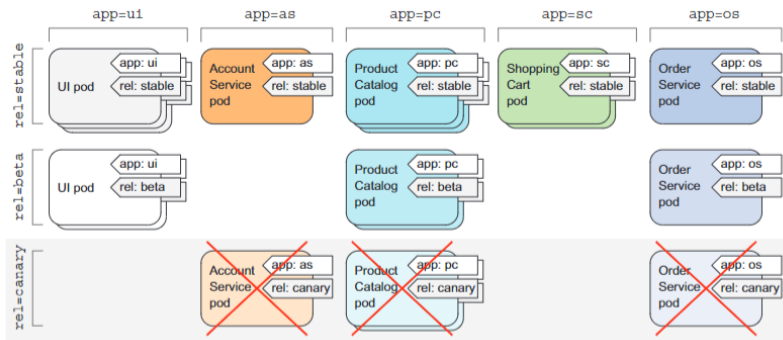
**Figure:** Listing by courtesy of Marko Luksa[1]

# Stopping and removing pods

Kubernetes sends a SIGTERM, waits 30 seconds, then SIGKILL.

```
$ kubectl delete po kubia-gp  
$ kubectl delete po -l creation_method=manual  
$ kubectl delete po -l rel=canary
```

# Stopping and removing pods cont.



**Figure:** Fig. by courtesy of Marko Luksa[1]

# Outline

Recap

Pods

**Services**

Deployments

# Services in Kubernetes


## Motivation

- ▶ We need a way to connect to pods from the outside.
- ▶ Pods are short-lived, they come and go.
- ▶ Clients should not know the IP's of pods.
- ▶ Scaling means multiple pods can provide the same service.

## How do they work?

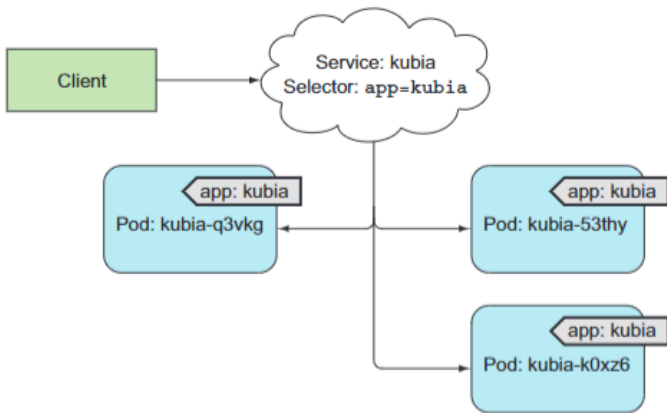
- ▶ A Service is a resource you create to make a single, constant point of entry to pods.
- ▶ Clients can now find frontend service, and frontend can find backend service.

## Services in Kubernetes cont.



**Figure:** Fig. by courtesy of Marko Luksa[1]

## Services in Kubernetes cont.



**Figure:** Fig. by courtesy of Marko Luksa[1]



## Services in Kubernetes cont.

An example of a YAML file for a service. Our service accepts connections port 80 and route each connection to port 8080 to a pod matching *app=kubia*.

```
apiVersion: v1
kind: Service
metadata:
  name: kubia
spec:
  ports:
  - port: 80
    targetPort: 8080
  selector:
    app: kubia
```

The port this service  
will be available on

The container port the  
service will forward to

All pods with the app=kubia  
label will be part of this service.

**Figure:** Listing by courtesy of Marko Luksa[1]

## Services in Kubernetes cont.

When created, we check to see if the service is running

```
$ kubectl get svc
NAME: kubia
CLUSTER-IP: 10.111.249.153
EXTERNAL-IP: <none>
PORT(S): 80/TCP
AGE: 6m
```

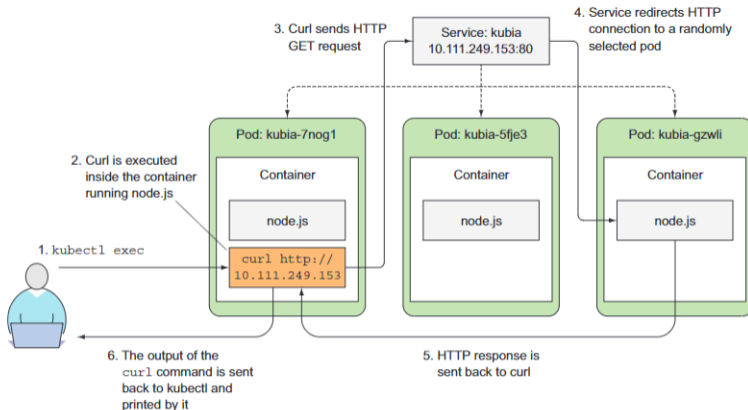
## Testing the service from the inside

With *kubectl exec* command we can access the service from inside  
Replace with your target pod and cluster IP

```
$ kubectl exec kubia-7nog1 -- curl -s \\  
http://10.111.249.153  
You've hit kubia-gzwli
```

The *kubectl exec* behaves similarly to *ssh*.

# Testing the service from the inside cont.



**Figure:** Listing by courtesy of Marko Luksa[1]


## Discovering our service via DNS

- ▶ Kubernetes comes with a *kube-dns* for free
- ▶ Each service gets a DNS entry in the internal DNS server
- ▶ Client pods can access the service by its FQDN.

```
$ kubectl exec -it kubia-3inly
root@kubia-3inly:/#
$ curl http://kubia.default.svc.cluster.local
Youve hit kubia-5asi2
$ curl http://kubia.default
You've hit kubia-3inly
$ curl http://kubia
You've hit kubia-8awf3
```

## Exposing services to external clients

- ▶ By setting the service type to NodePort (open up a port on the node itself)
- ▶ By setting the service type to LoadBalancer (deciated loud-balancer, AWS)
- ▶ By creating an Ingress resource (OSI level 7 resource)



**Figure:** Fig. by courtesy of Marko Luksa[1]

## Using a NodePort service

- ▶ For a NodePort service, each cluster node opens a port on the node itself and redirects traffic received on that port to the underlying service.
- ▶ This allows external traffic to our service.
- ▶ Can be configured with a YAML file.

```
apiVersion: v1
kind: Service
metadata:
  name: kubia-nodeport
spec:
  type: NodePort
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30123
  selector:
    app: kubia
```

Set the service  
type to NodePort.

This is the port of the  
service's internal cluster IP.

This is the target port  
of the backing pods.

The service will be accessible  
through port 30123 of each of  
your cluster nodes.

**Figure:** Listing by courtesy of Marko Luksa[1]

## Using a NodePort service cont.

We can check our service with *kubectl get svc* command

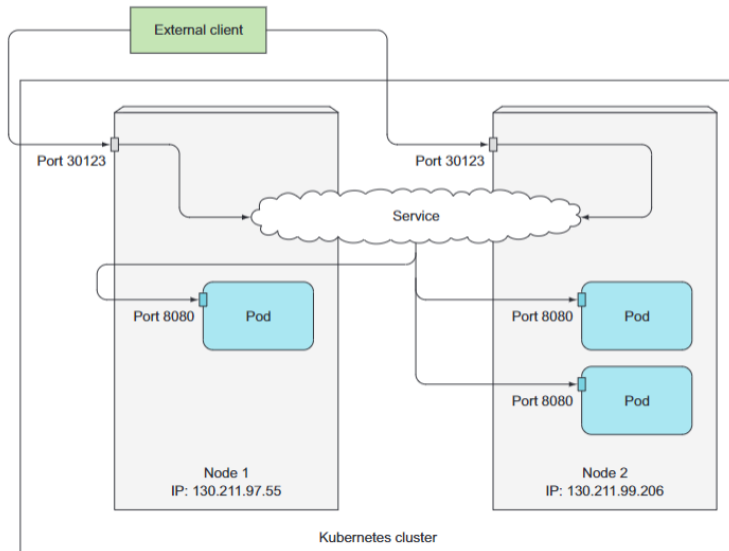
```
$ kubectl get svc kubia-nodeport
NAME: kubia-nodeport
CLUSTER-IP: 10.111.254.223
EXTERNAL-IP: <nodes>
PORT(S): 80:30123/TCP
AGE: 2m
```

The service is now accessible at the following addresses:

- ▶ 10.11.254.223:80
- ▶ <1st node's IP>:30123
- ▶ <2nd node's IP>:30123, and so on.



## Using a NodePort service cont.



## Using a LoadBalancer service

- ▶ A LoadBalancer service is the default way to expose a service to the Internet
- ▶ Each service gets its own IP
- ▶ Was not supported by minikube until recently
- ▶ Available by its EXTERNAL-IP

```
apiVersion: v1
kind: Service
metadata:
  name: kubia-loadbalancer
spec:
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: kubia
```

◀ This type of service obtains a load balancer from the infrastructure hosting the Kubernetes cluster.

**Figure:** Listing by courtesy of Marko Luksa[1]

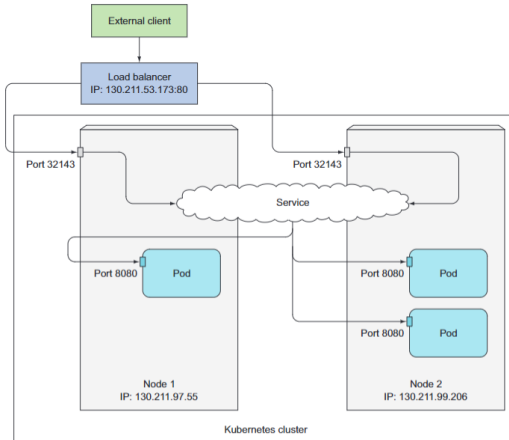
## Using a LoadBalancer service cont.

We can check our service with *kubectl get svc* command

```
$ kubectl get svc kubia-loadbalancer
NAME: kubia-loadbalancer
CLUSTER-IP: 10.111.241.153
EXTERNAL-IP: 130.211.53.173
PORT(S): 80:32143/TCP
AGE: 1m
```

The service is now accessible by its external IP:  
`curl http://130.211.53.173`

# Using a LoadBalancer service cont.



**Figure:** Fig. by courtesy of Marko Luksa[1]

## Using a headless service

Needed when client needs to connect to all pods

- ▶ A normal service just connects you to one pod
- ▶ A headless service will return the IP's of all annotated pods
- ▶ Can be accessed via DNS lookup
- ▶ Client libraries exist to ease communication with API

```
apiVersion: v1
kind: Service
metadata:
  name: kubia-headless
spec:
  clusterIP: None
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: kubia
```

← This makes the  
service headless.

**Figure:** Listing by courtesy of Marko Luksa[1]

## Using a headless service cont.

We can use the dnsutils image to verify a headless service

```
$ kubectl run dnsutils --image=tutum/dnsutils \\  
--generator=run-pod/v1 --command -- \\  
sleep infinity  
pod "dnsutils" created  
$ kubectl exec dnsutils nslookup kuba-headless  
...  
Name:      kuba-headless.default ...  
Address: 10.108.1.4  
Name:      kuba-headless.default ...  
Address: 10.108.2.5
```

# Outline

Recap

Pods

Services

**Deployments**

# Deployments in Kubernetes

What are deployments?

- ▶ A Deployment is a high-level resource
- ▶ It can deploy applications and update them declaratively
- ▶ Makes it easy to manage and make rolling-updates
- ▶ A Deployment is composed of a label selector, a replica count, and a pod template.

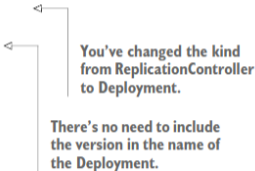


**Figure:** Fig. by courtesy of Marko Luksa[1]



## Deployments in Kubernetes cont.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubia
spec:
  replicas: 3
  template:
    metadata:
      name: kubia
      labels:
        app: kubia
    spec:
      containers:
        - image: luksa/kubia:v1
          name: nodejs
```



You've changed the kind from ReplicationController to Deployment.

There's no need to include the version in the name of the Deployment.

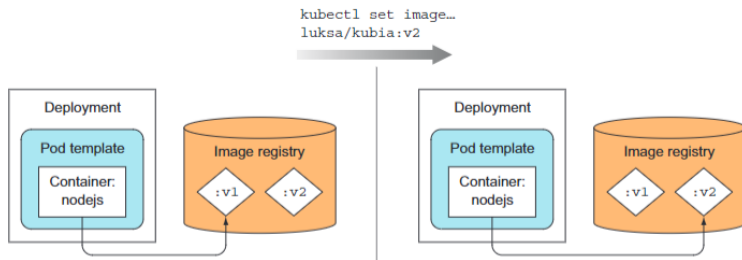
**Figure:** Listing by courtesy of Marko Luksa[1]

```
$ kubectl create -f kubia-deployment-v1.yaml
--record
$ kubectl rollout status deployment kubia
$ kubectl get po
```

## Rolling out updates


To roll-out an update, simply do

```
$ kubectl set image deployment kubia  
nodejs=luksa/kubia:v2
```



**Figure:** Fig. by courtesy of Marko Luksa[1]

# Rolling out updates cont.



**Figure:** Fig. by courtesy of Marko Luksa[1]

# Undoing a rollout


Deployments make it easy to roll back an update.

```
$ kubectl rollout undo deployment kubia
```

## View and using history

We can check the history of our deployments

```
$ kubectl rollout history deployment kubia  
$ kubectl rollout undo deployment kubia  
--to-revision=1
```



**Figure:** Fig. by courtesy of Marko Luksa[1]

## Control the rate of the rollout

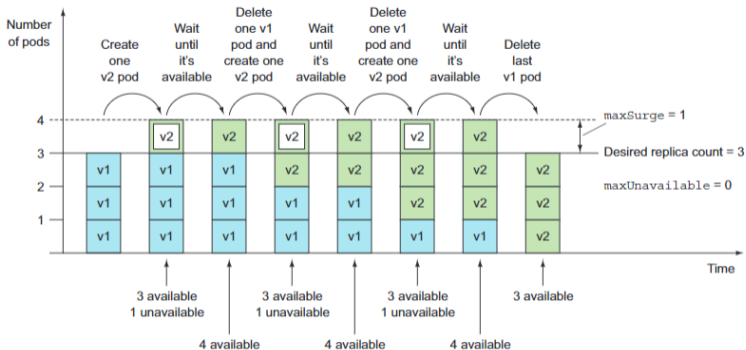
You can control the max surge and max unavailable pods in the Deployment manifest

```
spec:
  strategy:
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
    type: RollingUpdate
```

**Figure:** Listing by courtesy of Marko Luksa[1]


- ▶ **maxSurge:** Determines how many pods you allow to exist above the desired replica count. Default 25
- ▶ **maxUnavailable:** Determines how many pods can be unavailable relative to the desired replica count. Default 25

# Control the rate of the rollout cont.



**Figure:** Fig. by courtesy of Marko Luksa[1]

# Control the rate of the rollout cont.



**Figure:** Fig. by courtesy of Marko Luksa[1]



# Different ways to modify resources

There are several ways to modify a resource in Kubernetes

- ▶ `kubectl edit` - Opens a object's manifest in default editor
- ▶ `kubectl patch` - Changes individual properties of an object
- ▶ `kubectl apply` - Changes or creates an object from a YAML/JSON file
- ▶ `kubectl replace` - Replaces existing object from a YAML/JSON file
- ▶ `kubectl set image` - Changes a deployment's container image

# References I

- [1] Luksa, M. (2018). *Kubernetes in Action*. Manning Publications Co.